

# Uncertainty Quantification with Split Conformal Prediction

Thomas Davies

Edinburgh AI Society Advanced Workshop

November 4, 2024



- Machine learning appears hugely effective across many different problems
- However we don't know how certain we are about the predictions we make
  - A self-driving car needs to know when it's uncertain about an object
  - Medical diagnosis systems must express uncertainty in critical decisions
- Deep learning is notorious for being overly confident<sup>12</sup>
- Can we quantify uncertainty with minimal assumptions?

---

<sup>1</sup>Guo et al. "On Calibration of Modern Neural Networks"

<sup>2</sup>Begoli et al. "The need for uncertainty quantification in machine-assisted medical decision making"

# What is a Confidence Interval?

- Informally a confidence interval is an interval which is expected to typically contain the parameter being estimated.
- The confidence level (e.g. 95%) tells us how often the interval contains the true value
- If we repeat the experiment many times, 95% of the intervals will contain the true value
- A wider interval gives you more “confidence” but gives you less precision about the location of the parameter

# The Goal

- Given any supervised machine learning algorithm that maps paired training examples to a function:

$$(X_1, Y_1), \dots, (X_n, Y_n) \mapsto \hat{f}_{1:n} \quad (1)$$

- For a new data point  $X_{n+1}$ , we would like to construct a 95% “confidence set” for  $Y_{n+1}$
- This is a (random) set  $C(X_{n+1})$  where we have that

$$\mathbb{P}(Y_{n+1} \in C(X_{n+1})) = 0.95$$

- Without any assumptions on our learning algorithm, or underlying model this seems impossible!

# The Classical MLE Approach

- If we know the underlying distribution family, we can use Maximum Likelihood Estimation
- For example, with linear regression we assume:

$$Y = X\beta + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2)$$

- The MLE gives us asymptotically valid confidence intervals:

$$\hat{Y} \pm z_{1-\alpha/2} \hat{\sigma}$$

- But what if our distributional assumptions are wrong?
  - Non-normal errors
  - Non constant variance
  - Model misspecification

# No Asymptotic Results Needed!

These results also need infinite amounts of data for the confidence interval to be valid



To Infinity and Beyond!

We can do better!

# Motivation: Ranks of Random Variables

- Consider some i.i.d. real valued random variables  $Z_1, \dots, Z_{n+1}$
- What is the probability that  $Z_{n+1}$  is the  $k$ -th largest?
- By symmetry (exchangeability), all the orderings are equally likely
- Therefore:

$$\mathbb{P}(\text{Rank}(Z_{n+1}) = k) = \frac{1}{n+1} \quad (2)$$

- This holds regardless of the distribution of  $Z_i$ !

# A Note on Exchangeability

- Actually, we didn't need the i.i.d. assumption
- We only need exchangeability for any permutation  $\pi$ :

$$(Z_1, \dots, Z_{n+1}) \stackrel{d}{=} (Z_{\pi(1)}, \dots, Z_{\pi(n+1)})$$

- This is much weaker than i.i.d.
- Example: Drawing without replacement from an urn
  - Not independent (draws affect future probabilities)
  - But exchangeable (order doesn't matter)




# Building Confidence Intervals from Ranks

- If we have  $n + 1$  exchangeable random variables, we know their ranks are uniform
- We can use this to build confidence intervals!
- For a  $(1 - \alpha)$  confidence interval:
  - We want  $Z_{n+1}$  to be between the  $\lceil \frac{\alpha}{2}(n + 1) \rceil$ -th and  $\lfloor (1 - \frac{\alpha}{2})(n + 1) \rfloor$ -th order statistics
- This works regardless of the underlying distribution!

# Visualizing the Interval

- Assuming the variables are ordered as  $Z_1, Z_2, \dots, Z_n$  we can visualise this as.



# An Example

- With  $n = 99$  calibration points and  $\alpha = 0.05$ :
  - Lower bound: 3rd smallest value
  - Upper bound: 97th smallest value
  - Probability new point falls in this interval  $= 1 - \alpha = 0.95$
- This gives us a concrete way to build prediction intervals!
- And we still maintain our distribution-free guarantee

# Building Exchangeable Variables

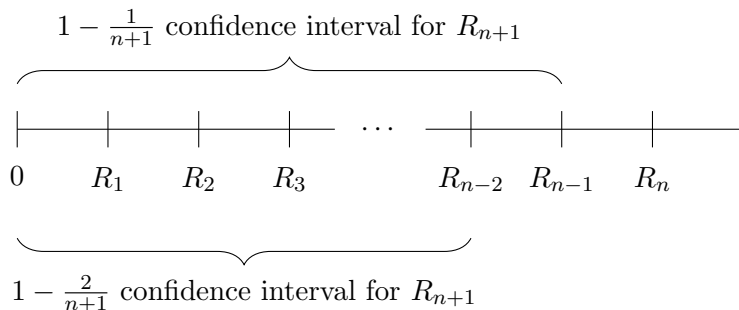
- Let's apply this to machine learning predictions by trying to build some exchangeable random variables - let's look at regression first.
- We split our training pairs into “pure training data”  $\mathcal{I}_1$  and “conformal calibrating data”  $\mathcal{I}_2$
- Let's train our model on  $\mathcal{I}_1$  to get  $\hat{f}(x)$  and then we can evaluate our model on  $\mathcal{I}_2$  to generate

$$R_i = |\hat{f}(X_i) - Y_i| \text{ for } (X_i, Y_i) \in \mathcal{I}_2$$

- These  $R_i$  are now exchangeable!!! So we can build a confidence interval for our next prediction using the ranking technique as before.

# Visualizing the Residuals

- We know that the residuals must be positive so we only need to worry about one side of the interval.
- Assuming the residuals are ordered as  $0, R_1, R_2, \dots, R_n$  we can visualise this as.



# The Split Conformal Prediction Algorithm

- For a new point  $X_{n+1}$ , compute  $R_{n+1}(y) = |\hat{f}(X_{n+1}) - y|$  for candidate values  $y$
- Our confidence set  $C(X_{n+1})$  is all values of  $y$  where  $R_{n+1}(y)$  is "not too large" compared to the calibration scores
- Specifically:  $y \in C(X_{n+1})$  if  $R_{n+1}(y)$  is smaller than the  $\lceil (1 - \alpha)(n + 1) \rceil$ -th largest calibration score
- This gives us valid  $(1 - \alpha)$  coverage by the rank arguments above!

# A Toy Example

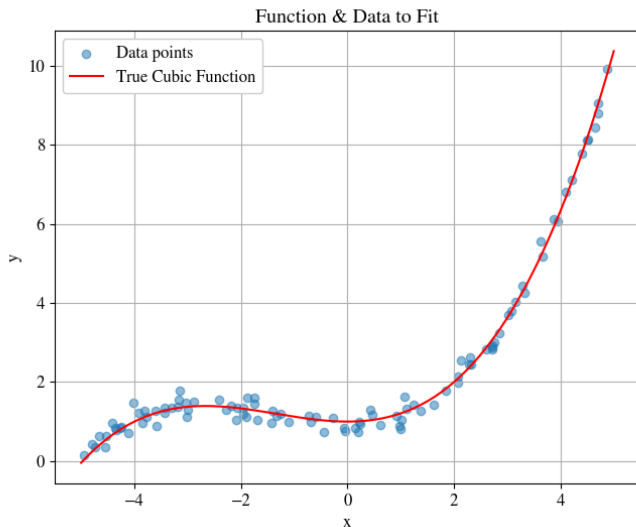
- Let's run a synthetic experiment to demonstrate conformal prediction
- We'll generate data from:
  - $X$  drawn from  $\text{Uniform}(-5, 5)$
  - $Y$  follows a cubic relationship with noise:

$$Y = 1 + \frac{1}{6}X^2 + \frac{1}{24}X^3 + \epsilon$$

where  $\epsilon \sim \mathcal{N}(0, 0.2)$

- This gives us ground truth to test drive our prediction intervals

# Visualizing the Data

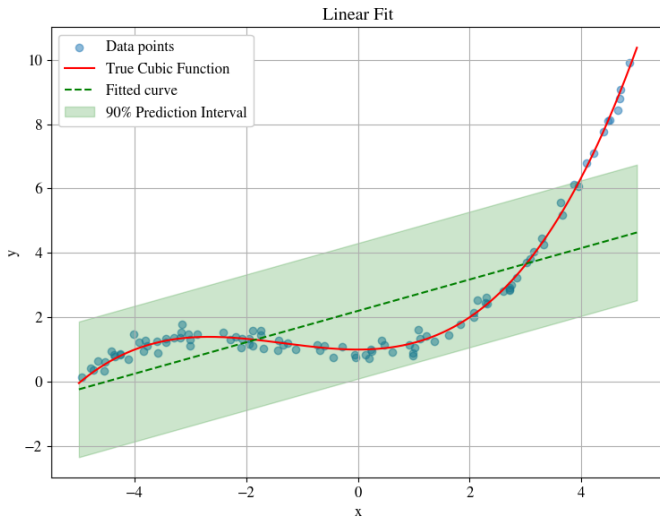




# Linear Model with Conformal Intervals

- Let's fit a linear model to this data - deliberately underfit!
- We should see that our conformal confidence interval gives a reasonably large interval
- The green region shows the 90% prediction interval
- Notice how:
  - The intervals are valid despite model misspecification
  - The intervals are very wide - we can't predict well so have high uncertainty


# Linear Fit with Conformal Intervals



# Linear Model with Conformal Intervals

- Let's now fit a cubic model to this data - this should give a much better fit
- We should see a small prediction interval

# Cubic Fit with Conformal Intervals



# Confidence Interval Size vs Model Complexity

Polynomial Degree	CI Width (lower is better)
0 (Constant)	5.03
1 (Linear)	3.82
2 (Quadratic)	2.40
3 (Cubic)	0.552
4 (Quartic)	0.583
5 (Quintic)	0.601

- Notice how the confidence interval width:
  - Decreases as we move from underfitting (degree 0) to good fit (degree 3)
  - Starts increasing again with higher degrees due to overfitting
- This shows that we can indeed use this to do model selection!

All of the conformal prediction was 9 lines of code!

```
def get_conformal_predictor(x, y, degree=3):  
    x_train, x_calib, y_train, y_calib = train_test_split(  
        x, y, test_size=0.5, random_state=42  
    )  
  
    coeffs = np.polyfit(x_train, y_train, degree)  
  
    y_pred_calib = np.polyval(coeffs, x_calib)  
    residuals = np.abs(y_calib - y_pred_calib)  
  
    def confinterval(x_new, alpha=0.9):  
        y_pred = np.polyval(coeffs, x_new)  
        q = np.quantile(residuals, alpha)  
        return y_pred - q, y_pred + q  
  
    return coeffs, confinterval
```

# What's Next?

- We can generalise our  $R_i$ s to be measures of a general “conformity score”  $V(x, y)$  and then

$$R_i = V(X_i, Y_i) \left( = |\hat{f}(X_i) - Y_i| \text{ in this talk} \right)$$

- All of our results still hold - there is nothing special about the absolute value. We can use this to build  $x$  adaptive predictive intervals (this is called Locally-Weighted Conformal Inference)<sup>3</sup>.
- Full Conformal Regression
- Leave one covariate out (LOCO) Conformal Regression

---

<sup>3</sup>Lei et al. “Distribution-Free Predictive Inference For Regression”

# Locally Adaptive Conformal Inference

Setup is the same as before, however now the noise ( $\epsilon$ ) has non constant variance.

